

# ClickHouse meets LLMs

Vector Search at Scale without a “Vector DB”

Jason Wong (Senior Support Engineer)



# ClickHouse is not “just OLAP”

it can be a **high-performance vector retrieval engine** in LLM systems.



# Everyone talks about LLMs. The real bottleneck is **retrieval at scale**.

- LLM magic is backed by RAG
- A successful RAG needs to...
  - Maintain high quality of embedding
  - Minimal retrieval latency
  - Holding an optimized storage
- Which product support all these????



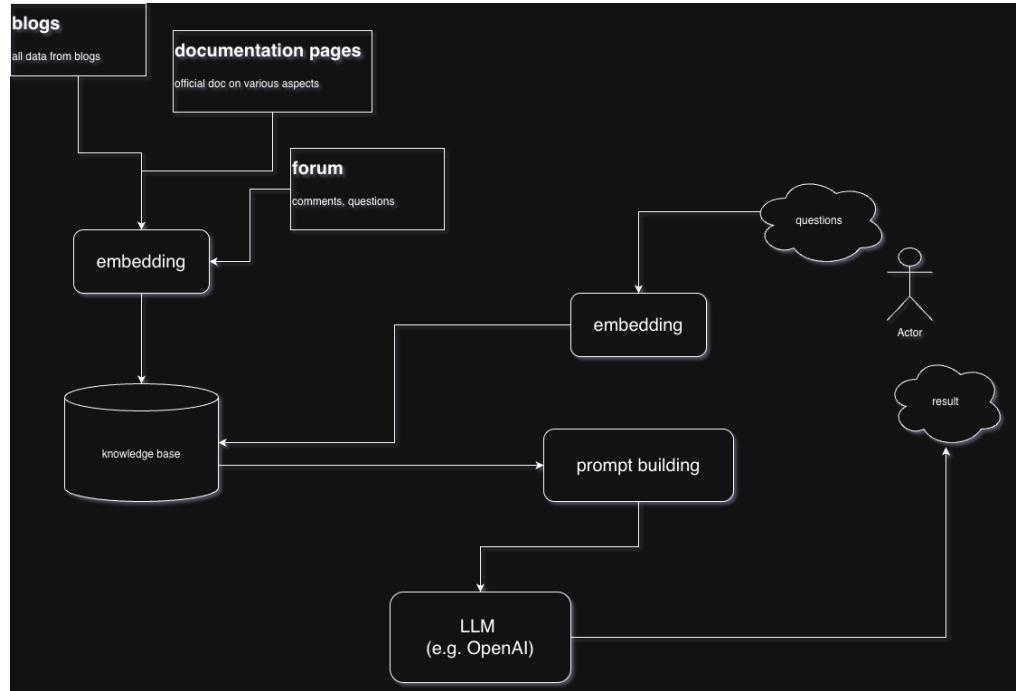
LLM: is a sophisticated type of artificial intelligence (AI) program that is trained on a massive amount of text data to understand, process, and generate human-like language

RAG: is an artificial intelligence (AI) technique that enhances the capabilities of large language models (LLMs) by connecting them to external knowledge sources



# Typical LLM Architecture

- Vector DB often introduced as yet another system
- Operational overhead
- Cost + data duplication



# Where ClickHouse Fits

- 1st class support for vectors (aka Arrays)
- Super fast ingestion and retrieval of data to provide real-time response
- ANN support (HNSW at the moment)
- Metadata filtering + vector search together

ANN: a computational model in artificial intelligence (AI) and machine learning (ML) that is inspired by the structure and function of the human brain's biological neural networks

HNSW: refers to the use of the Hierarchical Navigable Small World (HNSW) algorithm for performing Approximate Nearest Neighbor (ANN) search, a key technology in modern AI and machine learning applications



# Where ClickHouse Fits

- ClickHouse schema for the vector column is simply Arrays of `Float32`
- No special or dedicated data type needed - removing unnecessary type conversions for ingestion and retrieval

```
CREATE TABLE playground.docs
(
    doc_id        UInt64,
    chunk_id      UInt32,
    content       String,
    embedding     Array(Float32),
    source        LowCardinality(String),
    tenant_id     UInt32,
    ts            DateTime
)
ENGINE = MergeTree
ORDER BY (tenant_id, doc_id, chunk_id);
```



# Where ClickHouse Fits

- ClickHouse can filter meta-data which is not vector based (e.g. tenant\_id and source), not all vector-db support
- It also run KNN distance functions based on a vector column
- the BEST of both worlds

```
SELECT
    content,
    cosineDistance(embedding, [
        -0.0066272463, -0.07342478, -0.1604181, 0.031773847,
        -0.048222084, -0.023915742, 0.038059756, -0.017205033,
        -0.017487692, 0
    ]) AS score
FROM playground.docs
WHERE
    tenant_id = 88
    AND source = 'meetup_workshop'
    AND ts >= now() - INTERVAL 30 DAY
ORDER BY score asc
LIMIT 5;
```



# Embedding Generation

## Knowledge-base and user questions

- Language: Python 3
- Library / Module
  - sentence-transformers (Hugging Face)
  - clickhouse-connect (ClickHouse)



# Knowledge base Retrievals

- No SECRET at all, just a normal Query using ClickHouse built-in functions



# Takeaways

- One system (OLAP + Vector DB)
- Cost efficiency
  - columnar store
  - ZSTD compression
- Queries supporting meta-data and vector filtering



# About ME(me)

- Senior Support Engineer (ClickHouse)
- ElasticSearch Expert (full-text search, inverted-index...)
- Rust / Go / Vue.js Developer



LinkedIn



Demo Source code



# Thank you!

